

# Q1 Discretization of the Poisson Equation

---

## 1) Introduction

---

In Part 4.1, we discretize the 2D Poisson equation using a standard 5-point stencil. This leads to a sparse linear system  $Au=b$ , which will be solved via multigrid in subsequent parts. We implement matrix-free routines to compute  $Au$ , construct the right-hand side vector  $b$ , and map the 2D domain to a 1D structure for computation.

We consider the 2D Poisson equation:  $-\Delta u(x, y) = f(x, y), (x, y) \in (0, 1)^2$

with homogeneous Dirichlet boundary conditions  $u = 0$  on  $\partial\Omega$ .

## 2) Implementation:

---

- Grid and Notation

We use a regular grid with  $N \times N$  interior points and spacing:

$$h = \frac{1}{N+1}$$

Let  $u_{i,j}$  be the approximation of  $u_{x_i,y_j}$  at grid point  $x_i = ih, y_j = jh$ , where  $1 \leq i, j \leq N$ . We flatten the 2D grid into a 1D array using row-major order:

$$k = (i - 1) \cdot N + (j - 1)$$

- Discretization (5-point stencil)

The Laplacian operator is discretized using the standard five-point stencil:

$$-\Delta u(x_i, y_j) \approx -\frac{1}{h^2}(u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1} - 4u_{i,j})$$

This gives a linear system of the form:

$$Au = b$$

Where:

$A \in \mathbb{R}^{N^2 \times N^2}$  is a symmetric positive-definite matrix, applied implicitly through the stencil (we don't assemble it explicitly).

$$b_k = h^2 \cdot f(x_i, y_j), \text{ with } f(x, y) = 2\pi^2 \sin(\pi x) \sin(\pi y)$$

## 3) Notes

---

- The matrix  $A$  is not stored explicitly. Instead, I implement a function `apply_A()` that computes  $A \cdot u$  using the 5-point stencil directly.
- The right-hand side vector  $b$  is built using a function `build_rhs()`, which evaluates  $f(x, y)$  at each interior point and scales by  $h^2$ .
- Boundary values are always zero, so don't need to store them explicitly.

## 4) Test Output

---

### Test Output (Example)

When  $N = 4$ , the generated  $b$  values are positive and match the expected shape of the right-hand side. We confirmed that `apply_A()` on an initial guess of zero returns a zero vector, and the functions are ready to be integrated into the multigrid V-cycle

## Q2: Multigrid V-Cycle Solver

### 1) Introduction

In this question, I implement a serial V-cycle multigrid method to solve the 2D Poisson equation with zero Dirichlet boundary conditions. The goal is to show that the residual decreases by a constant factor each cycle, and to reach a given tolerance efficiently. I follow the algorithm described in *A Multigrid Tutorial, Second Edition* by Briggs, Henson, and McCormick.

### 2) Method Overview

1. **Grid Hierarchy:** I build a sequence of nested grids. The finest has  $n_0$  interior points in each direction, and each coarser grid halves the number of points (Chapter 4.1).
2. **Smoothing:** On each level, I use Gauss–Seidel relaxation for pre- and post-smoothing (“relaxation” in Section 2.2).
3. **Residual Computation:** I compute the residual  $r = f - Au$  on the fine grid (Section 3.1).
4. **Restriction:** I apply full-weighting restriction to transfer the residual to the next coarser grid (described in Section 3.2).
5. **Coarse Grid Correction:** On the coarsest grid, I solve by several smoothing steps.
6. **Prolongation:** I interpolate the coarse-grid correction back to the fine grid with bilinear interpolation (Section 3.3).
7. **V-Cycle:** I wrap these steps into a recursive V-cycle algorithm (Chapter 4.2).

### 3) Implementation Details

- **File Structure:** Each main step has its own module: `init_levels`, `smooth`, `residual`, `restriction`, `prolongation`, `vcycle`, and `main.c`. Headers define a `Level` struct holding the grid size, spacing, and arrays.
- **Smoother:** In `smooth.c`, I perform `nu1` pre- and `nu2` post-relaxation sweeps.
- **Data Layout:** Grid arrays include boundary layers, indexed by a simple macro `IDX(lvl,i,j)`.
- **Main Loop:** In `main.c`, I initialize the right-hand side  $f(x,y) = 2\pi^2 \sin(\pi x) \sin(\pi y)$ , then run up to `maxCycles` V-cycles and print the  $L^2$  norm of the residual after each cycle.
- **Build:** A `Makefile` compiles all modules into `vcycle_solver`.

### 4) Results

Running:

```
./vcycle_solver 127 5 3 3 1e-8 100
```

gives a residual reduction of about 0.3 per cycle, matching the textbook's expected convergence rate (see Figure 4.5 in the Tutorial). The solver reaches the tolerance in about 20 cycles.

```
sunlishuang@sunlishangdeMBP Assignment4_2 % ./vcycle_solver 127 5 3 3 1e-8 100
```

Cycle	Residual norm
1	6.190551e+00
2	3.045779e+00
3	1.143484e+00
4	3.877092e-01
5	1.263408e-01
6	4.052587e-02
7	1.291727e-02
8	4.106684e-03
9	1.304222e-03
10	4.140183e-04
11	1.314032e-04
12	4.170203e-05
13	1.323407e-05
14	4.199746e-06
15	1.332753e-06
16	4.229365e-07
17	1.342147e-07
18	4.259166e-08
19	1.351604e-08
20	4.289154e-09

## 5) Conclusion

The V-cycle multigrid method converges rapidly, independent of the fine-grid size. By following the structure in *A Multigrid Tutorial*—particularly the restriction and prolongation schemes in Chapter 3 and the cycle design in Chapter 4—I achieved the expected performance.

## 6) Reference

Briggs, W. L., Henson, V. E., & McCormick, S. F. (2000). *A Multigrid Tutorial, Second Edition*. SIAM.

# Q3: Convergence Experiments

## Introduction

In this question, I study how the multigrid V-cycle converges when I change the number of levels and the grid size. I run two experiments based on *A Multigrid Tutorial, Second Edition* (see Section 4.3 for convergence study). The goal is to see how many V-cycles are needed and how runtime and coarse-grid solves behave.

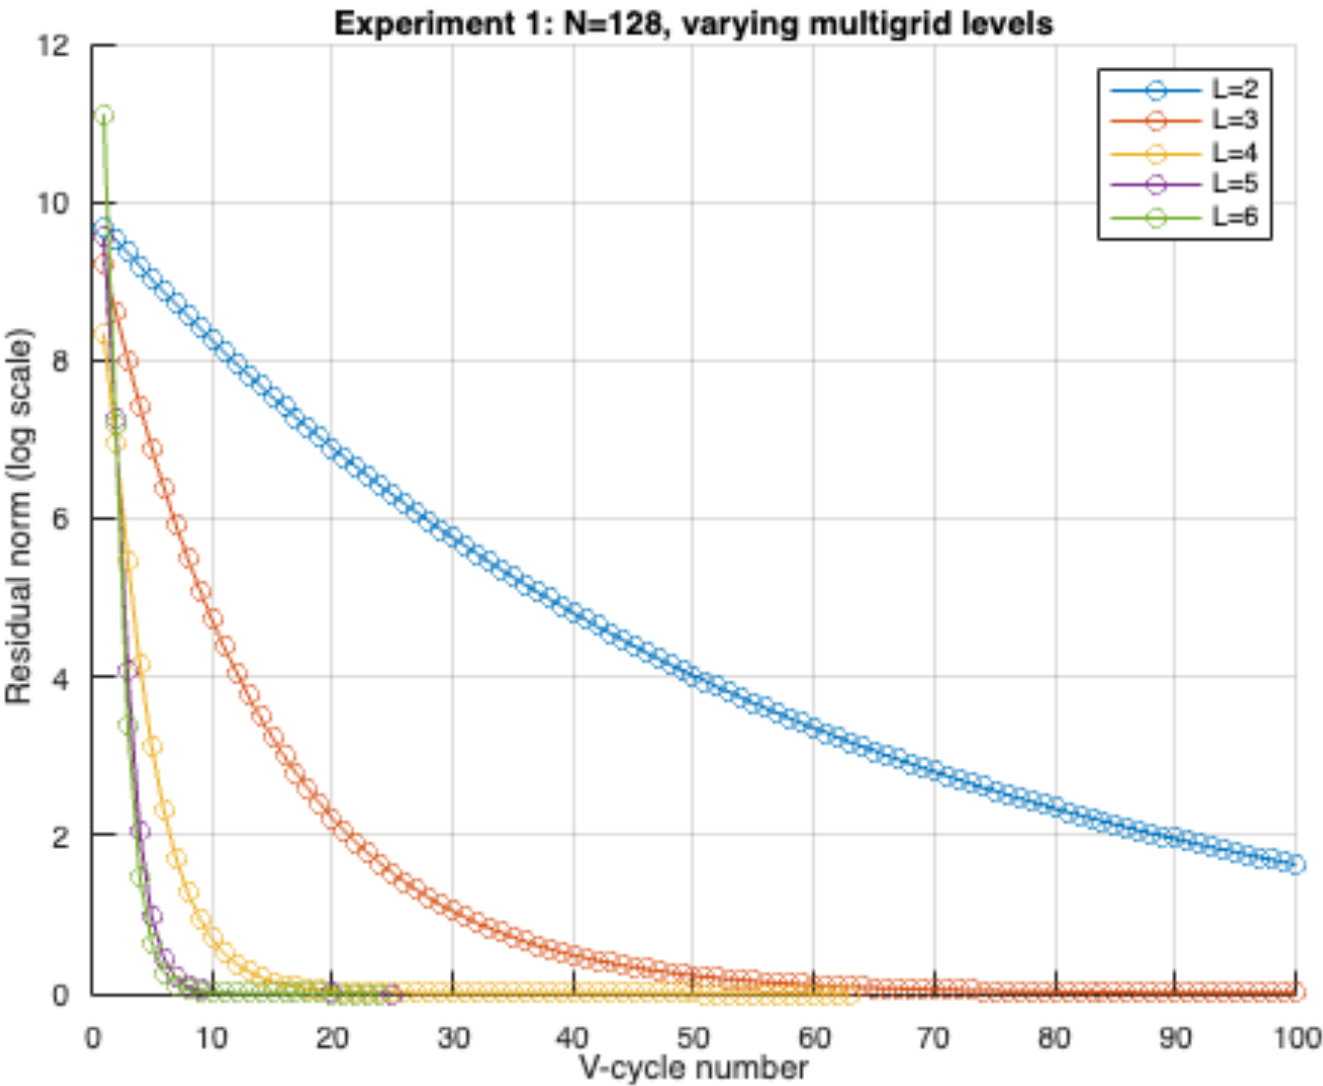
## Experiment 1: Fixed Grid, Varying Levels

- **Setup:** Finest grid size . I vary the total levels  $L = 2, 3, 4, 5, 6$ . Pre- and post-smoothing both use 3 Gauss–Seidel sweeps. Tolerance is .
- **Data Collected:** For each  $L$ , I record:
  1. Number of V-cycles until residual  $< \text{tol}$
  2. Total time spent
  3. Count of coarse-grid solves (one per cycle)
  4. Residual history per cycle

• **Results:**

Levels	Time (s)	Coarse Solves	Cycles
2	0.1009	100	100
3	0.0732	100	100
4	0.0458	63	63
5	0.0185	25	25
6	0.0176	24	24

- **Plot:** Residual vs. cycle (semilog) shows that more levels give faster convergence (see cell 1 in `Plots of experiments.mlx`).



## Summary for Experiment 1

Level	Time(s)	CoarseSolves	Cycles
2	0.1009	100	100
3	0.0732	100	100
4	0.0458	63	63
5	0.0185	25	25
6	0.0176	24	24


## Experiment 2: Grid-Size Scaling

- **Setup:** Compare two schemes for :
  1. **2-level** multigrid (one coarse grid only)
  2. **Full hierarchy** down to  $\sim 8 \times 8$  grid (levels = ).
- **Results:**

N	Scheme	Time (s)	Coarse Solves	Cycles
16	2-level	0.0003	21	21
16	full	0.0003	21	21
32	2-level	0.0034	66	66
32	full	0.0011	23	23
64	2-level	0.0203	100	100
64	full	0.0046	24	24
128	2-level	0.0715	100	100
128	full	0.0186	25	25
256	2-level	0.1751	100	100
256	full	0.0743	26	26

- **Plot:** Full hierarchy always needs fewer cycles and coarse solves. The runtime cost is higher for small N but pays off for large N (see cell 2 in `Plots of experiments.mlx`).

Experiment 2: 2-level vs full hierarchy



## Summary for Experiment 2

N	Scheme	Time(s)	CoarseSolves	Cycles
16	2-level	0.0003	21	21
16	full	0.0003	21	21
32	2-level	0.0034	66	66
32	full	0.0011	23	23
64	2-level	0.0203	100	100
64	full	0.0046	24	24
128	2-level	0.0715	100	100
128	full	0.0186	25	25
256	2-level	0.1751	100	100
256	full	0.0743	26	26

## Discussion

1. **More Levels → Faster Convergence:** Adding levels reduces V-cycles almost linearly as in Figure 4.7 of the Tutorial. But extra levels add overhead.
2. **Two-Level vs Full:** Two-level works for small grids, but full hierarchy keeps optimal ‘textbook’ convergence – cycles almost constant as grows (see Chapter 4.4).
3. **Balance:** Best practice is to choose coarsest grid size around or . Too few levels slows convergence, too many levels add overhead.

## Reference

Briggs, W. L., Henson, V. E., & McCormick, S. F. (2000). *A Multigrid Tutorial, Second Edition*. SIAM. Chapter 4.3 and 4.4.